


# Learning one transition at a time

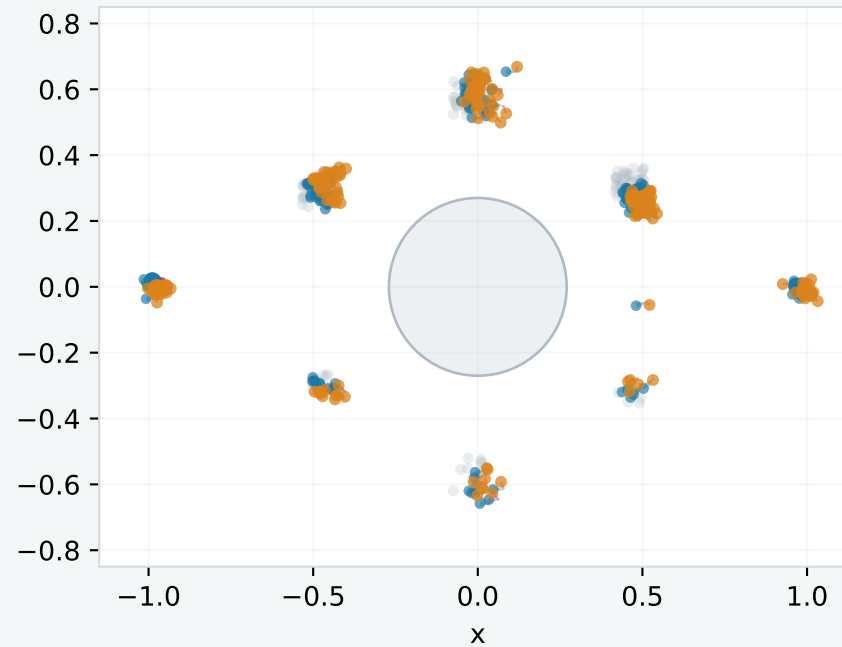
Shared-state encoder • held-out interpolation scene 1 • class 0 (80% upper-route preference)

Top: independent transitions at five saved times. Bottom: midpoint detail; all three panels use identical axes.

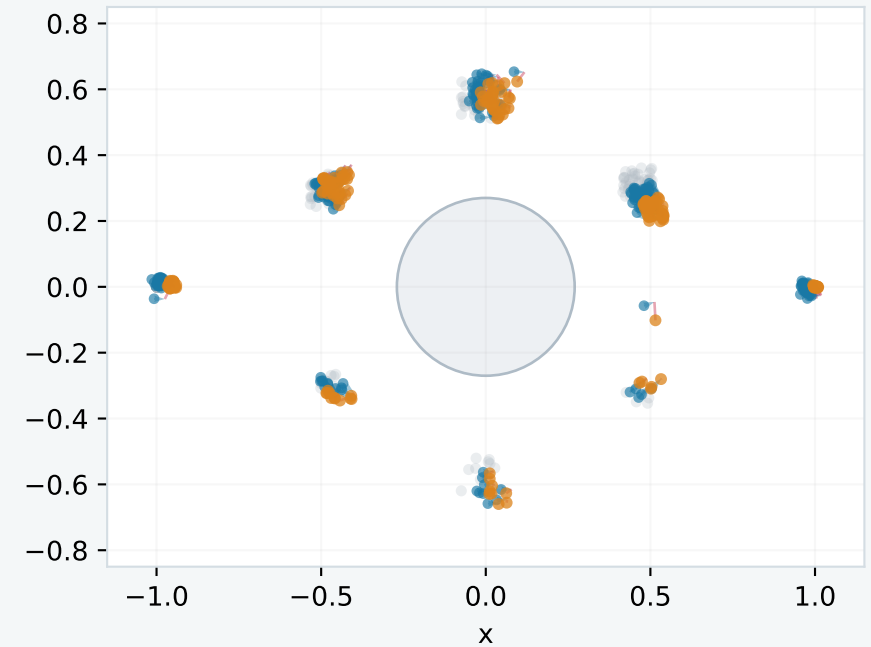
Real transitions



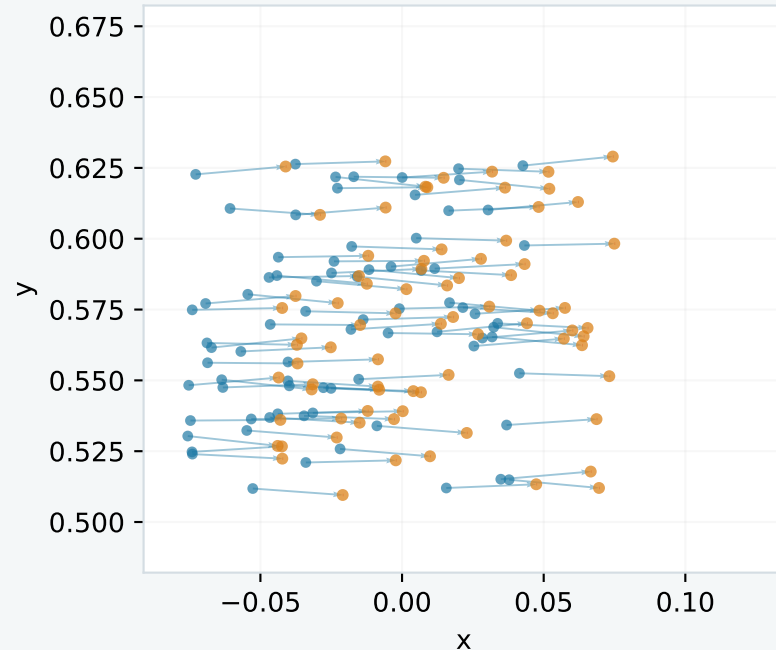
G1 / G2 / G3



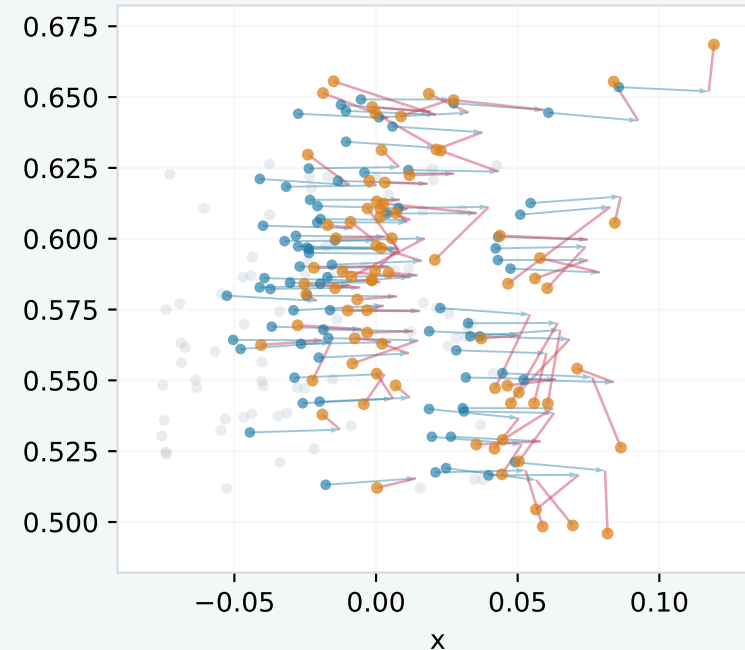
G1 / G2 → E → G3



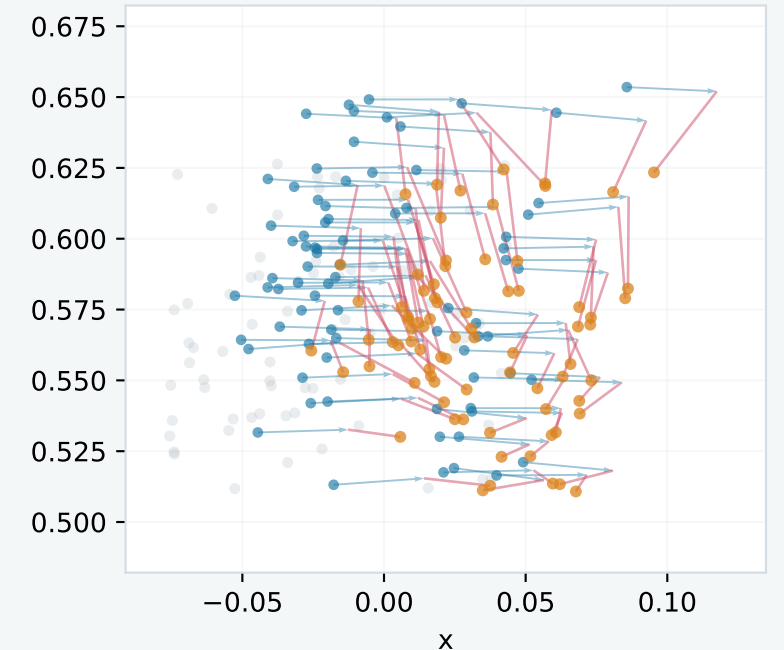
Midpoint, upper branch • residual 0.0000



Midpoint, upper branch • residual 0.0176



Midpoint, upper branch • residual 0.0226



Blue point + arrow:  $st$  and  $at$     Orange point: generated  $st+1$     Red line: disagreement with  $st + at$

Original joint SW1: 0.08565 (14.6% better than prior leader). Extrapolation coverage: 0.72%. This is a toy result, not solved dynamics.